

Final Project of Machine Learning with Spark

Author - Rohit Rokde

```
In [1]: import findspark
findspark.init('/u/rrokde/spark-2.4.4-bin-hadoop2.7')
from pyspark import SparkConf, SparkContext
from pyspark.sql import SparkSession

conf = SparkConf().setMaster("local").setAppName("Project Assignment - ML with Spark")
sc = SparkContext(conf=conf)
sc.setLogLevel("ERROR")
spark = SparkSession.builder.getOrCreate()
```

```
In [2]: # 1. Data Loading

data = spark.read.load('/u/rrokde/mls/default of credit card clients1.csv',
                        format='csv', header=False, inferSchema=True, sep=',')
```

```

In [3]: # 2. Data Transformation

#We need to remove the garbage values from the first row.
from itertools import islice
data = data.rdd.mapPartitionsWithIndex(lambda idx, it: islice(it, 1, None) if
(idx == 0 ) else it).toDF()
firstRow = data.head(1)

#Second row contains the column names.
columnNames = []
for i in range(len(firstRow[0])):
    columnNames.append( str(firstRow[0][i]).replace(' ', '_') )
data = data.rdd.mapPartitionsWithIndex(lambda idx, it: islice(it, 1, None) if
(idx == 0 ) else it).toDF()

#Next we tidy up the dataframe by setting the right column names and casting c
olumn data to the right type
for i in range(len(firstRow[0])):
    data = data.withColumnRenamed("_c" + str(i), columnNames[i])
    data = data.withColumn(columnNames[i],data[columnNames[i]].cast('double'))
    # Source: https://forums.databricks.com/questions/9147/how-to-infer-csv-sc
hema-default-all-columns-like-s.html

#Now we optimize our data for machine learning. That is extracting features an
d then normalizing it.
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.feature import StringIndexer
from pyspark.ml.feature import VectorIndexer
from pyspark.ml.feature import MinMaxScaler

#Code taken from 1st assignment
df = data
list_of_columns = df.columns
list_of_columns.pop() #Remove the target column from feature vector
assembler = VectorAssembler().setInputCols(list_of_columns).setOutputCol("feat
ures")
transformed = assembler.transform(df)
scaler = MinMaxScaler(inputCol="features",outputCol="scaledFeatures")
scalerModel = scaler.fit(transformed.select("features"))
scaledData = scalerModel.transform(transformed)

#Remove unnecessary columns. We have all the data in vectorised column - scale
dFeatures
for col in scaledData.columns:
    if( (col == 'default_payment_next_month') or (col == 'scaledFeatures')):
        pass
    else:
        scaledData = scaledData.drop(col)

print("Dataframe structure used for training and evaluation:-")
scaledData.printSchema()

```

Dataframe structure used for training and evaluation:-

```
root
|-- default_payment_next_month: double (nullable = true)
|-- scaledFeatures: vector (nullable = true)
```

```
In [4]: # 3. Model Learning

# Split the data into training and test sets
(trainingData, testData) = scaledData.randomSplit([0.7, 0.3])

#Source: https://spark.apache.org/docs/latest/ml-classification-regression.html#binomial-logistic-regression
from pyspark.ml.classification import DecisionTreeClassifier
from pyspark.ml import Pipeline

# Train a DecisionTree model.
dt = DecisionTreeClassifier(labelCol="default_payment_next_month", featuresCol="scaledFeatures")

# Chain indexers and tree in a Pipeline
#pipeline = Pipeline(stages=[LabelIndexer, featureIndexer, dt])
pipeline = Pipeline(stages=[dt])

# Train model. This also runs the indexers.
model = pipeline.fit(trainingData)
```

```
In [5]: # 4. Model evaluation

# Make predictions.
predictions = model.transform(testData)

from pyspark.ml.evaluation import MulticlassClassificationEvaluator
# Select (prediction, true label) and compute test error
evaluator = MulticlassClassificationEvaluator(
    labelCol="default_payment_next_month", predictionCol="prediction", metricName="accuracy")

accuracy = evaluator.evaluate(predictions)
print("DecisionTreeClassifier accuracy = ", str(format( (accuracy)*100, '.2f')) + "%")

DecisionTreeClassifier accuracy = 81.69%
```

Accuracy of the model is around 82%

```
In [6]: # Select example rows to display.  
print("Some values to verify our results:-")  
predictions.select("prediction", "default_payment_next_month").show(10)
```

Some values to verify our results:-

```
+-----+-----+  
|prediction|default_payment_next_month|  
+-----+-----+  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
|      0.0|                        0.0|  
+-----+-----+
```

only showing top 10 rows